

# Lecture # 1 : Introduction

September 12, 2023

# What is HPC ?

## *High-Performance Computing (HPC)*

Mixture of

1. Specialized (parallel) hardware
  2. Adapted (parallel) algorithms
  3. *Ad hoc* (parallel) programming techniques
- + Specially trained programmers !

# Course Topics

1. Introduction
2. Distributed programming using MPI
3. Multi-thread programming with OpenMP
4. Vectorization
5. Understanding the hardware and exploiting its full capacity
6. Recurring algorithmic themes/problems

Next year: GPUs/accelerators

# Fundamental Role of Numerical Simulations

## Traditional scientific method

1. Formulation of a theory based on observations
2. Design of experiments to (in)validate the theory
3. Compare experimental results with observations

Iterate the process until convergence.

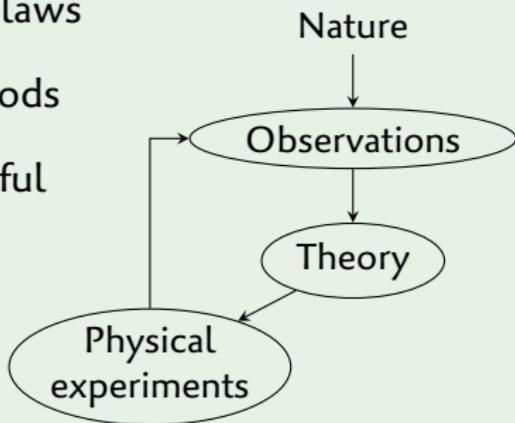
## But experimenting can be too...

- ▶ Difficult (*Wind Tunnel* containing an Airbus A380?)
- ▶ Expensive (*crash test, ...*)
- ▶ Slow (climate evolution, galaxy dynamics, ...)
- ▶ Dangerous (clinical trials, ☣, ☢, ☣—, ...)

# Numerical Simulations

Using computers to simulate and analyze a physical phenomenon

- ▶ Starting from physical laws
- ▶ Using numerical methods
- ▶ with ever more powerful computers

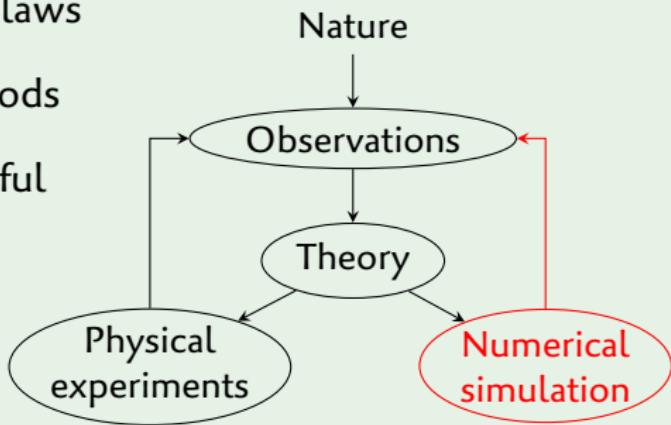


~ 3rd pillar of science?

# Numerical Simulations

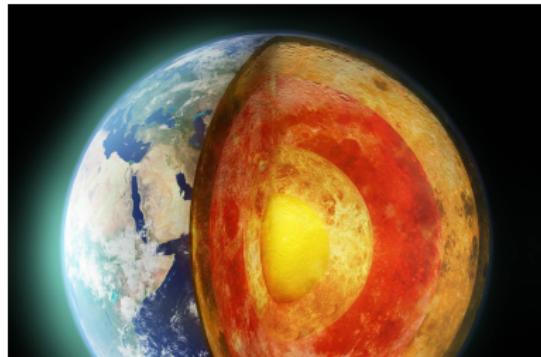
Using computers to simulate and analyze a physical phenomenon

- ▶ Starting from physical laws
- ▶ Using numerical methods
- ▶ with ever more powerful computers



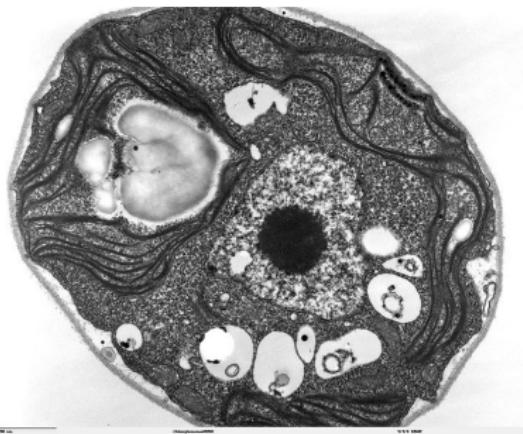
~~~ 3rd pillar of science?

## Illustration of the Problem



- ▶ surface = 510 millions  $\text{km}^2$ , volume  $\approx 10^{12} \text{ km}^3$
- ▶ One double per  $\text{km}^3 \rightsquigarrow 8 \text{ TB of memory}$
- ▶ One double per  $\text{m}^2 \rightsquigarrow 4 \text{ PB of memory}$

## Illustration of the Problem (continued)



- ▶ Human cell. About  $10^{14}$  atoms.
- ▶  $3 \times$  double  $(x,y,z)$  per atom  $\rightsquigarrow 2.4$  PB of memory

Not to mention...

- ▶  $10^{11}$ – $10^{12}$  stars in our galaxy
- ▶  $10^{11}$ – $10^{12}$  galaxies...

# Realization



- ▶ You know **a little** about how computers works
  - ▶ But **not completely**
  - ▶ Especially not **BIG** ones
- ▶ You have **not** been trained to :
  - ▶ **Care** about the **performance** of your code
  - ▶ Write **efficient** programs

# Realization



- ▶ You know **a little** about how computers works
  - ▶ But **not completely**
  - ▶ Especially not **BIG** ones
- ▶ You have **not** been trained to :
  - ▶ **Care** about the **performance** of your code
  - ▶ Write **efficient** programs
- ▶ **Time for a change**

# Computing Power

## performance measurement

Unit : FLOP (*Floating Point OPeration*), FLOP/s (or FLOPS)

- ▶ Giga ( $10^9 \approx 2^{30}$ )
- ▶ Tera ( $10^{12} \approx 2^{40}$ )
- ▶ Peta ( $10^{15} \approx 2^{50}$ )
- ▶ Exa ( $10^{18} \approx 2^{60}$ ) ...

- ▶ Machines for large computations = **parallel computers**
- ▶ Parallelism is an **inexorable trend**

# 1990's PC

The only one you really know how to program



1 core, 25Mhz, 16KB cache, 4MB RAM

# Contemporary Laptop

What you own, fortunately



4 cores, 4Ghz, 8MB cache, 16GB RAM (2 channels)

# Gamer's PC

For smart-asses who think AI will replace actual science



10–16 cores, 4Ghz, 16MB cache, 32GB RAM (2 channels)

# Compute Node

Now we're talking



2 CPUs, 32–128 cores, 2.5Ghz, 256GB RAM (2× 6–8 channels)

# Supercomputer

Rare and Expensive



1000–10000000 cores. Cost between  $10^6$  and  $10^9$  euros

# Description of an HPC machine: Turing

IBM BlueGene/Q (2012–2019) @ IDRIS (CNRS)



# Description of an HPC machine: Turing

IBM BlueGene/Q (2012–2019) @ IDRIS (CNRS)

6 racks, 98 304 cores ("small")  $\leadsto$  Sequoia (USA) has 1 572 864



# Description of an HPC machine: Turing

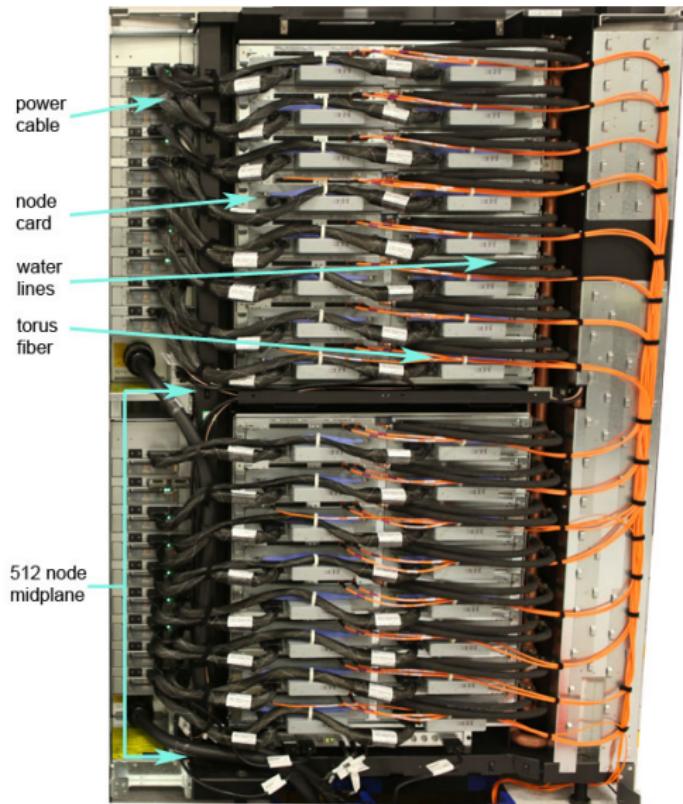
IBM BlueGene/Q (2012–2019) @ IDRIS (CNRS)

1 rack =

- ▶ 16 *node cards*
- ▶ + power feeds
- ▶ + *water cooling*
- ▶ + I/O nodes

Total =

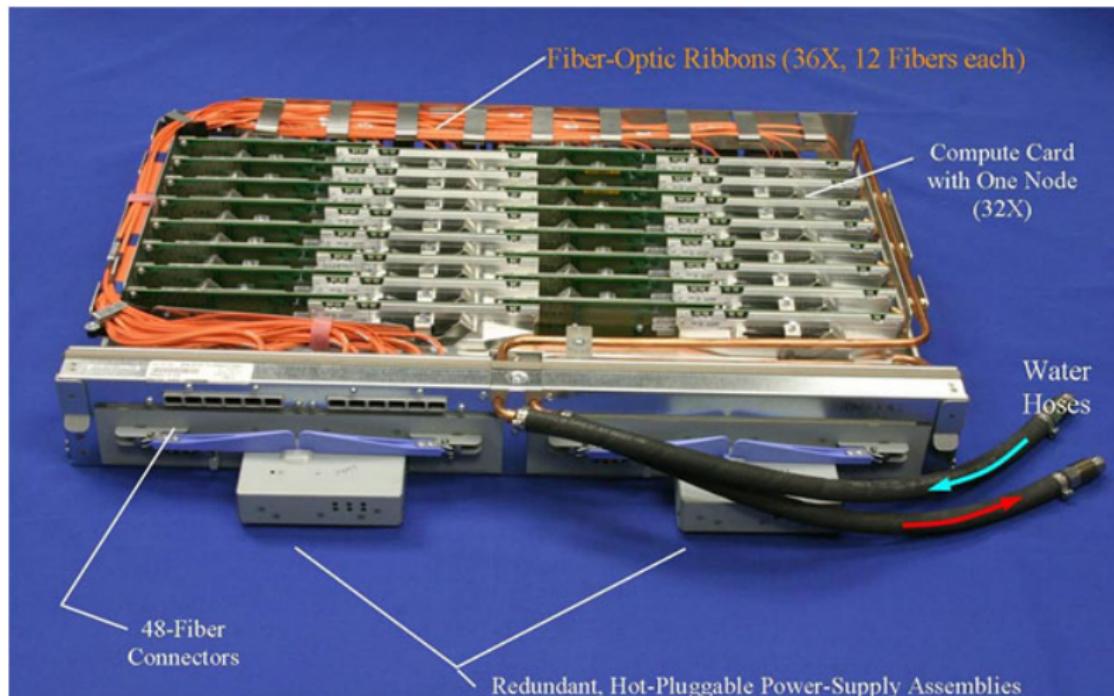
- ▶ 1024 nodes
- ▶ 16 384 cores
- ▶ 16 TB RAM



# Description of an HPC machine: Turing

IBM BlueGene/Q (2012–2019) @ IDRIS (CNRS)

1 node card = 32 compute cards



# Description of an HPC machine: Turing

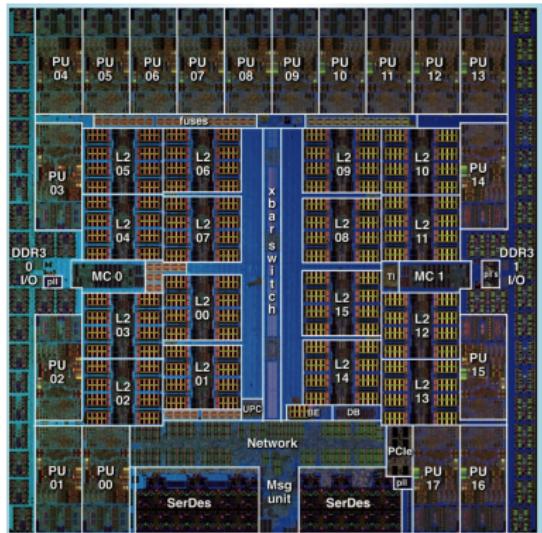
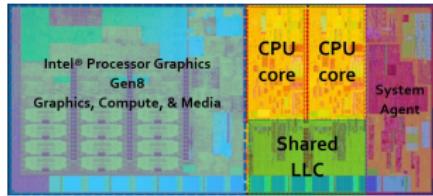
IBM BlueGene/Q (2012–2019) @ IDRIS (CNRS)

1 *compute card* =

- ▶ 1 PowerPC A2 CPU (16 cores @ 1.6Ghz, 32MB Cache)
- ▶ 16GB of ECC RAM (36 × 512MB chip)



# Not “Normal” computers



## A Few Examples (in France)

1. Turing (IDRIS / CNRS, 2012–2019), 1.2 PFLOPS
  - ▶ IBM BlueGene/Q ( $\approx$  20 millions €)
  - ▶ 6144 nodes, 98 304 cores, 98TB RAM
    - ▶ 16GB + 1 × IBM PowerPC A2, 16 cores @ 1.6 Ghz
    - ▶ 1 CPU = ??? € = 200 GFLOPS
  - ▶ Interconnect = “in-house” 5D-torus, 160 Gbit/s per node
  - ▶ 600kW (cold water cooling)

## A Few Examples (in France)

1. Turing (IDRIS / CNRS, 2012–2019), 1.2 PFLOPS
  - ▶ IBM BlueGene/Q ( $\approx$  20 millions €)
  - ▶ 6144 nodes, 98 304 cores, 98TB RAM
    - ▶ 16GB +  $1 \times$  IBM PowerPC A2, 16 cores @ 1.6 Ghz
    - ▶ 1 CPU = ??? € = 200 GFLOPS
  - ▶ Interconnect = “in-house” 5D-torus, 160 Gbit/s per node
  - ▶ 600kW (cold water cooling)
2. Jean-Zay (IDRIS / CNRS, 2019–), 28 PFLOPS
  - ▶ HPE SGI 8600 (25 millions €)
  - ▶ 1528 “CPU” nodes  $\rightarrow$  61 120 cores, 4.9 PFLOPS
    - ▶ 192GB +  $2 \times$  Intel Xeon Gold 6248, 20 cores, 2.5Ghz
    - ▶ 1CPU = 3000€ = 1.6 TFLOPS
  - ▶ Interconnect = 100Gbit/s OmniPath
  - ▶  $\approx$  1MW (warm water cooling)

## A Few Examples (in France)

1. Turing (IDRIS / CNRS, 2012–2019), 1.2 PFLOPS
  - ▶ IBM BlueGene/Q ( $\approx$  20 millions €)
  - ▶ 6144 nodes, 98 304 cores, 98TB RAM
    - ▶ 16GB +  $1 \times$  IBM PowerPC A2, 16 cores @ 1.6 Ghz
    - ▶ 1 CPU = ??? € = 200 GFLOPS
  - ▶ Interconnect = “in-house” 5D-torus, 160 Gbit/s per node
  - ▶ 600kW (cold water cooling)
2. Jean-Zay (IDRIS / CNRS, 2019–), 28 PFLOPS
  - ▶ HPE SGI 8600 (25 millions €)
  - ▶ 1528 “CPU” nodes  $\rightarrow$  61 120 cores, 4.9 PFLOPS
    - ▶ 192GB +  $2 \times$  Intel Xeon Gold 6248, 20 cores, 2.5Ghz
    - ▶ 1CPU = 3000€ = 1.6 TFLOPS
  - ▶ 662 “GPU” nodes  $\rightarrow$  199 040 cores, 21.5 PFLOPS
    - ▶ idem +  $4 \times$  NVIDIA V100 SMX2 16/32GB, 80 cores
    - ▶ 1 GPU = 8500€ = 8 TFLOPS
  - ▶ Interconnect = 100Gbit/s OmniPath
  - ▶  $\approx$  1MW (warm water cooling)

## More Examples (TOP 500)

1. #5 = Summit (DoE, USA, 2018–), 200 PFLOPS
  - ▶ IBM ( $\approx$  \$ 500 millions)
  - ▶ 4 608 nodes
    - ▶ 2× IBM Power9, 22 cores, 3Ghz
    - ▶ 6× NVIDIA V100 16GB
    - ▶ 512GB DDR4 + 96GB HBM + 1.6TB non-volatile memory
    - ▶ 1 node  $\approx$  120 000€ = 42 TFLOPS
  - 2.8PB RAM, 202 752 CPU cores, 2 211 840 GPU cores
  - ▶ Interconnect = 100Gb/s Infiniband
  - ▶ 13MW

## More Examples (TOP 500)

### 1. #5 = Summit (DoE, USA, 2018–), 200 PFLOPS

- ▶ IBM ( $\approx \$ 500$  millions)
- ▶ 4 608 nodes
  - ▶ 2 × IBM Power9, 22 cores, 3Ghz
  - ▶ 6 × NVIDIA V100 16GB
  - ▶ 512GB DDR4 + 96GB HBM + 1.6TB non-volatile memory
  - ▶ 1 node  $\approx 120\ 000\text{€} = 42$  TFLOPS
- 2.8PB RAM, 202 752 CPU cores, 2 211 840 GPU cores
- ▶ Interconnect = 100Gb/s Infiniband
- ▶ 13MW

### 2. #21 = Frontera (Univ. Texas, 2019–), 38.7 PFLOPS

- ▶ Dell ( $\approx \$60$  millions)
- ▶ 8008 nodes
  - ▶ 192GB + 2 × Intel Xeon Platinum 8280, 28 cores @ 2.7Ghz
  - ▶ 1 CPU = 10 000€ = 2.4 TFLOPS
- 1.5PB RAM, 448 448 cores
- ▶ Interconnect = 100Gbit/s Infiniband

## And of Course...

- ▶ #2 = Fugaku (RIKEN, Japon, 2021–), 550 PFLOPS
  - ▶ Fujitsu ( $\approx$  900 millions €)
  - ▶  $\geq 150\,000$  nodes
    - ▶ 32GB HBM2 + 1× Fujitsu A64FX (arm), 48 cores @ 2Ghz
    - ▶ 1 CPU = ??? € = 2.7 TFLOPS
  - 4.8PB RAM, 7 200 000 cores
  - ▶ Interconnect = “in-house” 3D torus

## And of Course...

- ▶ #2 = Fugaku (RIKEN, Japon, 2021–), 550 PFLOPS
  - ▶ Fujitsu ( $\approx$  900 millions €)
  - ▶  $\geq$  150 000 nodes
    - ▶ 32GB HBM2 + 1× Fujitsu A64FX (arm), 48 cores @ 2Ghz
    - ▶ 1 CPU = ??? € = 2.7 TFLOPS
  - 4.8PB RAM, 7 200 000 cores
  - ▶ Interconnect = “in-house” 3D torus
- ▶ #1 = Frontier (DOE, USE, 2022–), 1.1 EFLOPS
  - ▶ HPE Cray ( $\approx$  600 millions €)
  - ▶  $\geq$  9472 nodes
    - ▶ 512GB RAM + AMD Epyc 7A53, 64 cores @ 2Ghz
    - ▶ 4× AMD Radeon Instinct MI250X GPUs (128GB)
  - 606,208 CPU cores + 8,335,360 GPU cores + 9PB RAM
  - ▶ Interconnect = HPE Slingshot
  - ▶ 21MW

# Your (Near) Future...



- ▶ #9999 = One computer lab room,  $\approx 0.8$  TFLOPS
  - ▶ 16 nodes
    - ▶ 4GB RAM + 1 × Core i5 CPU (4 cores)
    - ▶ 1 CPU = 50 GFLOPS
  - 64GB RAM, 64 cores
  - ▶ Interconnect = (gigabit ?) ethernet

# Great Architectural Diversity

## How to get computing power ?

- ▶ As-powerful-as-possible CPUs (frontera)
- ▶ Tons of weak CPUs (IBM BlueGene)
- ▶ Hardware accelerators (GPU) (frontier, summit, jean-zay)
- ▶ CPU-GPU hybrids (fugaku)

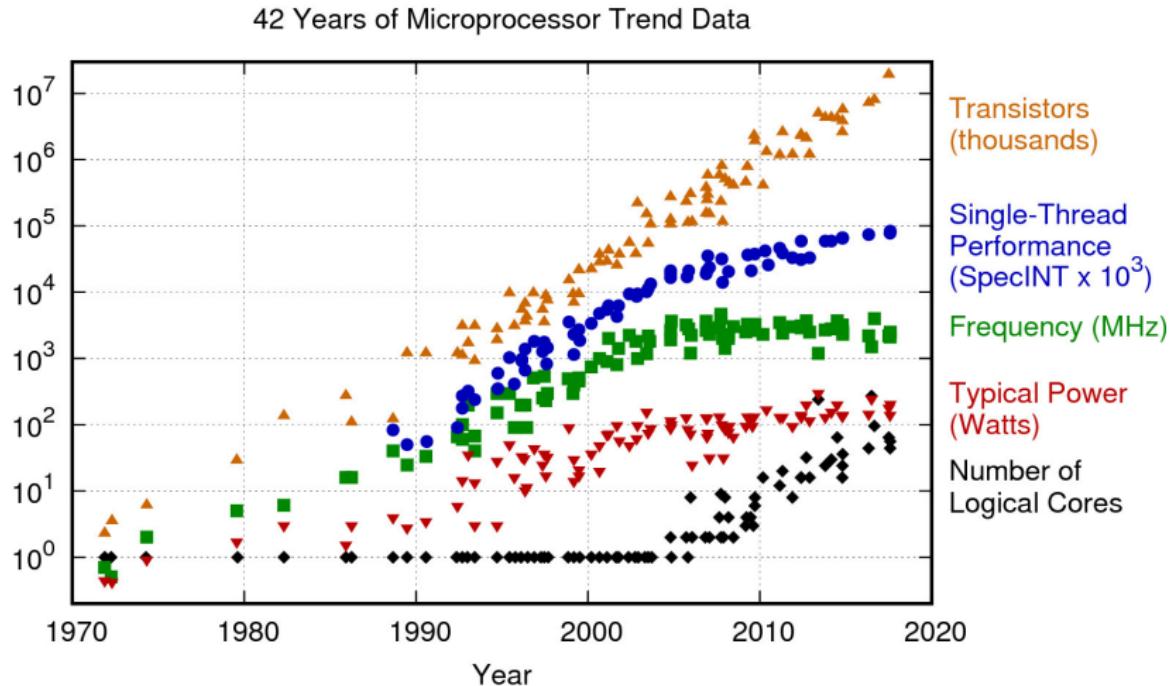
## Porting difficulties

Some codes are sometimes written for **ONE** target machine

# TOP500 History

| Machine               | year | nodes   | M/node | FLOPS  |
|-----------------------|------|---------|--------|--------|
| Frontier              | 2022 | 9 472   | 1T     | 1100P  |
| Fugaku                | 2020 | 158 976 | 32G    | 430P   |
| Summit                | 2018 | 4 608   | 608G   | 150P   |
| Sunway TaihuLight     | 2016 | 40 960  | 32G    | 93P    |
| Tiahne-2A             | 2013 | 16 000  | 88G    | 62P    |
| Titan                 | 2012 | 18 688  | 38G    | 18P    |
| Sequoia               | 2012 | 98 304  | 16G    | 17P    |
| K computer            | 2011 | 82 944  | 16G    | 11P    |
| Tiahne-1A             | 2010 | 7 168   | 32G    | 2 566T |
| Jaguar                | 2009 | 18 688  | 16G    | 1 941T |
| Roadrunner            | 2008 | 3 240   | 32G    | 1 105T |
| BlueGene/L            | 2004 | 106 496 | 512M   | 478T   |
| Earth Simulator       | 2002 | 640     | 16G    | 39T    |
| ASCI White            | 2000 | 512     | 16G    | 7226T  |
| ASCI Red              | 1997 | 4 649   | 256M   | 2379G  |
| CP-PACS/2048          | 1996 | 2 048   | 64M    | 368G   |
| Numerical Wind Tunnel | 1993 | 166     | 25M    | 124G   |
| CM-5                  | 1993 | 1 024   | 3M     | 60G    |

# "Moore's Law" and the End of "Dennard Scaling"

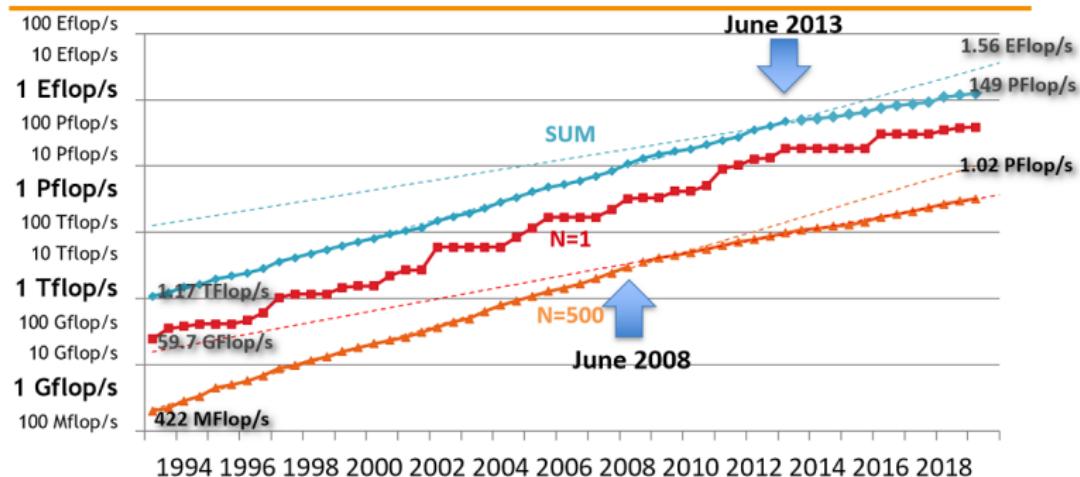


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2017 by K. Rupp

# "Moore's Law" and the End of "Dennard Scaling"

Consequence in the TOP500

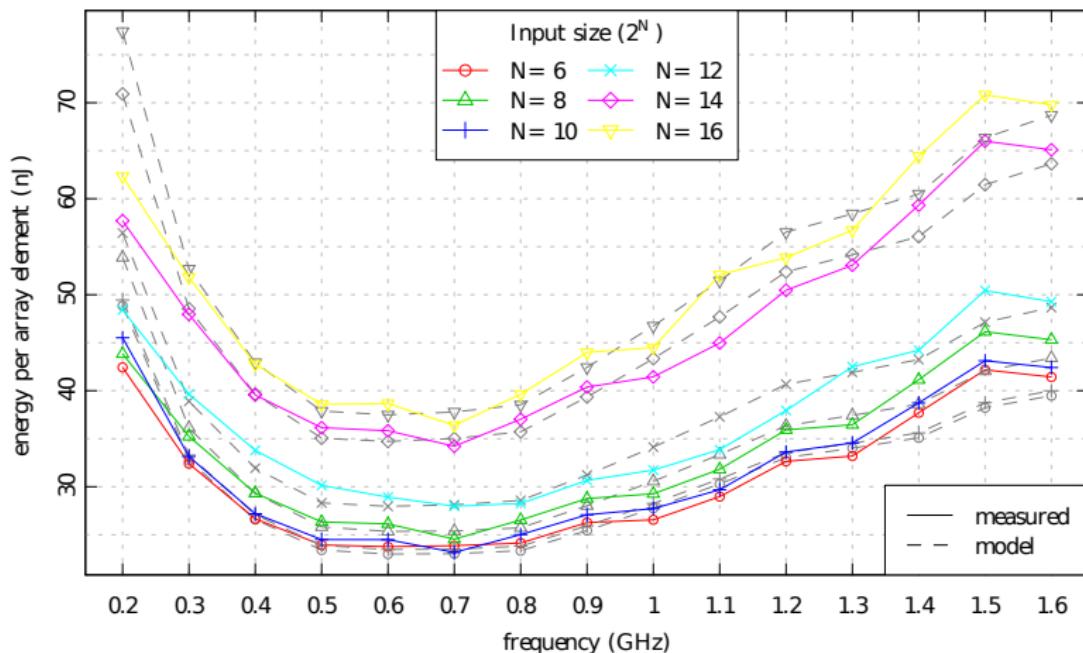
## PERFORMANCE DEVELOPMENT



# Improving the Energy Efficiency (FLOP/Watt)

~~ Reduce Frequency, increase # processors

CPU = Cortex A9 (ARM, smartphone). Computation = (part of the) FFT.



# Glossary

machine all components

cluster compute servers connected to a network

node an independent “computer”

rack frame or enclosure for mounting servers, switches...

SMP node with several processors

processor contains at least one core + cache, etc.

socket connects a processor to the motherboard

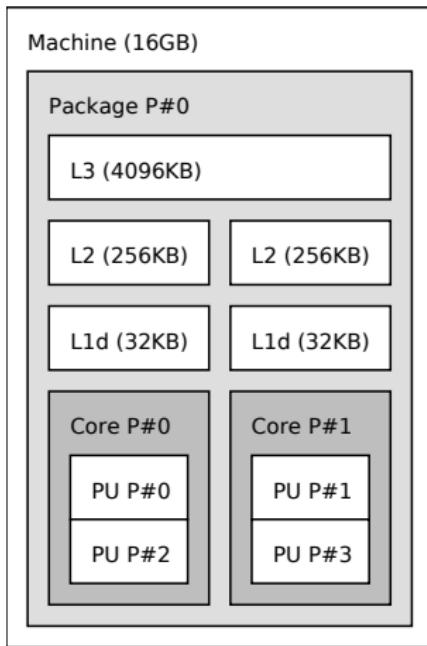
core circuit executing code autonomously

multicore processor containing several cores

hardware thread autonomous execution context in a core

SMT core that hosts several hardware threads

## Glossary – trouble with “*thread*”



`lstopo` output on a laptop

- ▶ hardware thread  $\neq$  software thread (OS)
- ▶ “physical core”  $\approx$  core
- ▶ “logical core”  $\approx$  hardware thread

# Parallel Machines – Classification

Flynn's taxonomy (1966) :

|                    |          | Data stream |             |
|--------------------|----------|-------------|-------------|
|                    |          | single      | multiple    |
| Instruction Stream | single   | <b>SISD</b> | <b>SIMD</b> |
|                    | multiple | <b>MISD</b> | <b>MIMD</b> |

# Parallel Machines – Classification

Flynn's taxonomy (1966) :

|                    |          | Data stream |             |
|--------------------|----------|-------------|-------------|
|                    |          | single      | multiple    |
| Instruction Stream | single   | <b>SISD</b> | <b>SIMD</b> |
|                    | multiple | <b>MISD</b> | <b>MIMD</b> |

- ▶ **SISD** Usual sequential machines
- ▶ **MISD** Do not really exist

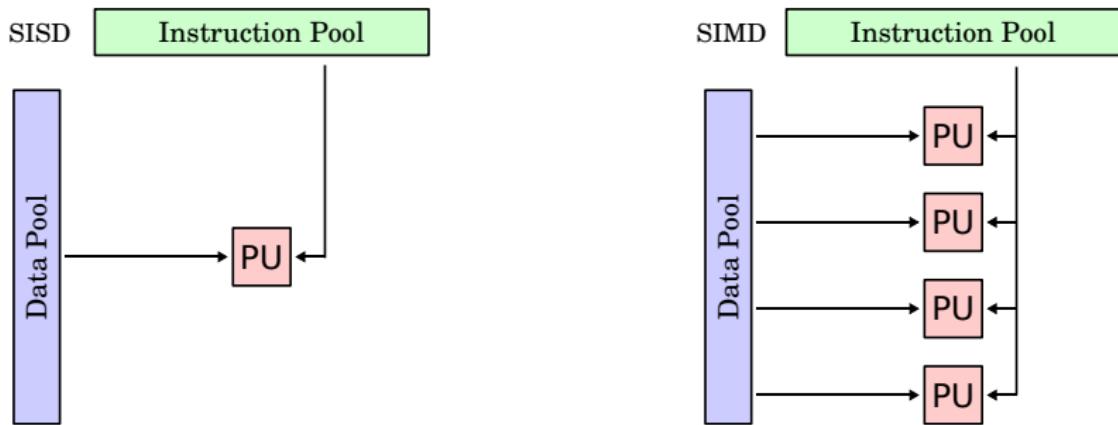
## Flynn's Taxonomy (continued)

### ► SIMD

Processing units simultaneously execute the same operation on their own data

- ▶ Synchronous operations
- ▶ Single centralized control unit
- ▶ Examples :
  - ▶ Vector machines from the 1970–90 (CRAY , NEC ...)
  - ▶ Vector instructions in contemporary CPUs (SSE, AVX2, AVX-512, NEON, ...)
  - ▶ Graphics Processing Units (GPUs)

# SIMD



(image: Wikipedia)

# Flynn's Taxonomy (continued)

## ► SIMD

Processing units simultaneously execute the same operation on their own data

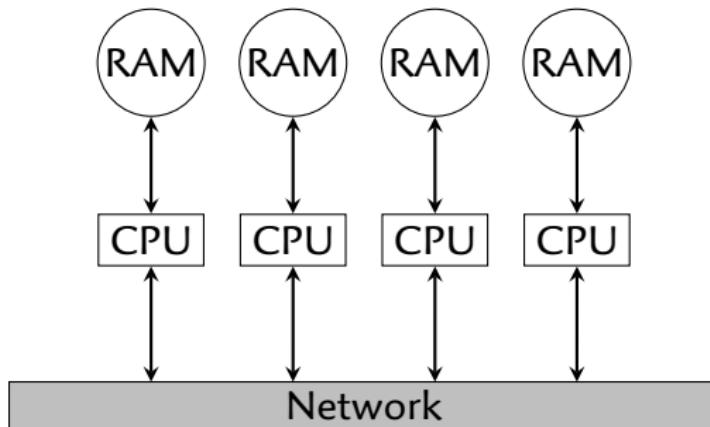
- ▶ Synchronous operations
- ▶ Single centralized control unit
- ▶ Examples :
  - ▶ Vector machines from the 1970–90 (CRAY , NEC ...)
  - ▶ Vector instructions in contemporary CPUs (SSE, AVX2, AVX-512, NEON, ...)
  - ▶ Graphics Processing Units (GPUs)

## ► MIMD

Independent processors can perform arbitrary operations on their own data simultaneously.

- ▶ Asynchronous
- ▶ Examples :
  - ▶ Multicore CPU
  - ▶ Clusters
  - ▶ All big HPC machines

# Distributed Memory



- ▶ Each processor has its own memory
- ▶ Communication = messages exchange on a network
- ▶ Network performance = **critical** aspect of the machine

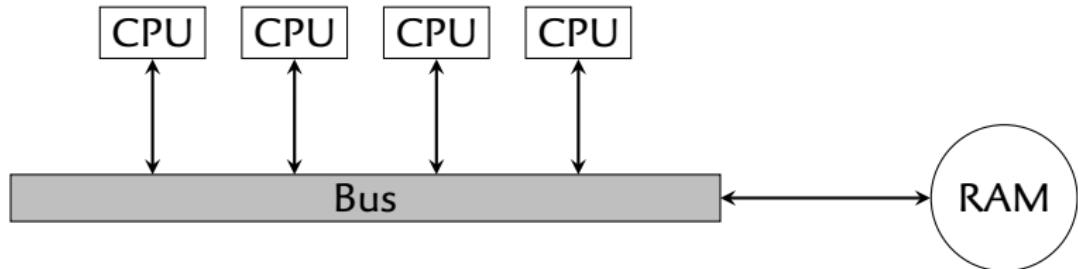
# Bad Idea

Tree-like Ethernet, Ordinary Switches



Bad performance, scalability problems

# Shared Memory



- ▶ CPUs connected to the memory through a bus
  - ▶ Each CPU may access the whole memory
  - ▶ Communications ↔ memory accesses
- 
- ▶ Beware of race conditions!

# Non-Uniform Memory Access

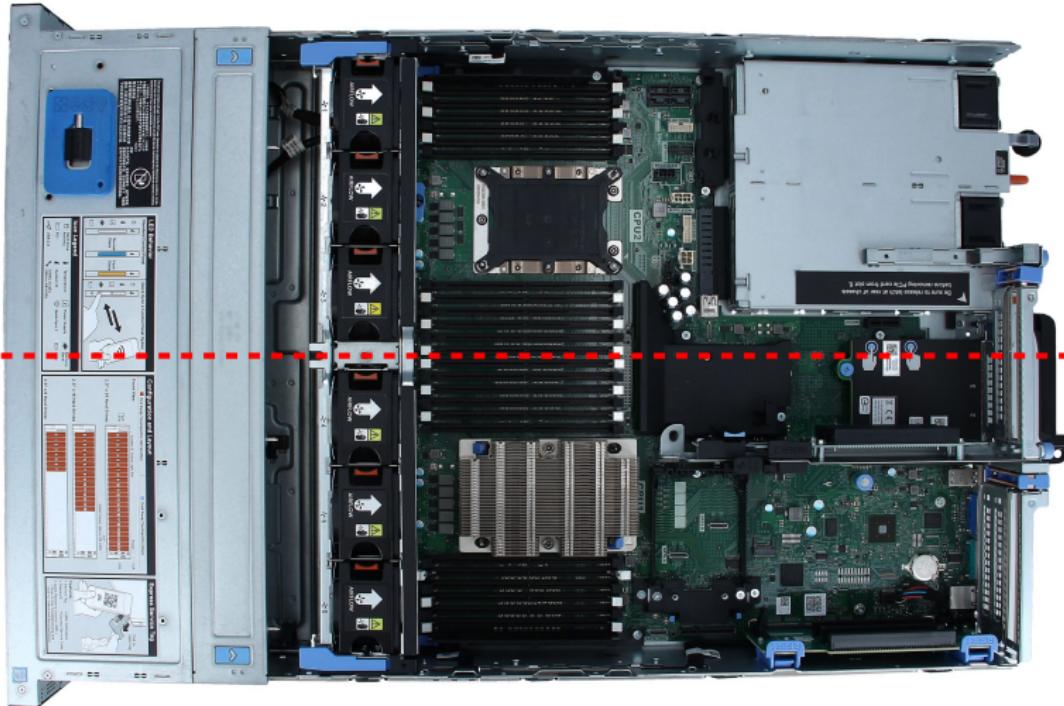
Shared memory: scalability problem

More CPUS  $\rightsquigarrow$  contention on the bus

## Solution

- ▶ Simulate a shared memory with a distributed memory
  - ▶ Each CPU “owns” and control a share of the memory
    - ▶ *Directly* accesses its own share
  - ▶ Accesses the rest by communicating with other CPUs
- ⇒ *ad hoc* network, very fast.

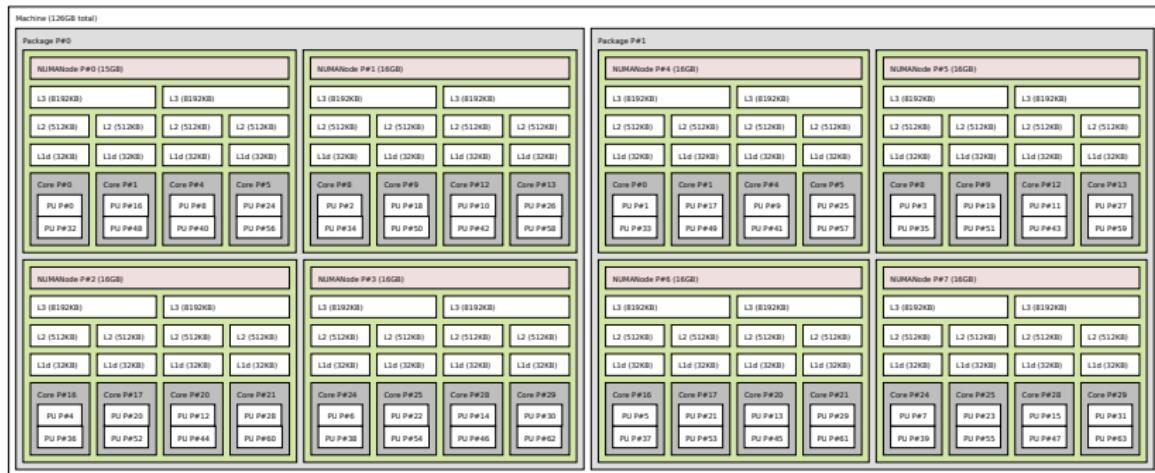
# Non-Uniform Memory Access



# Non-Uniform Memory Access

## Real Example

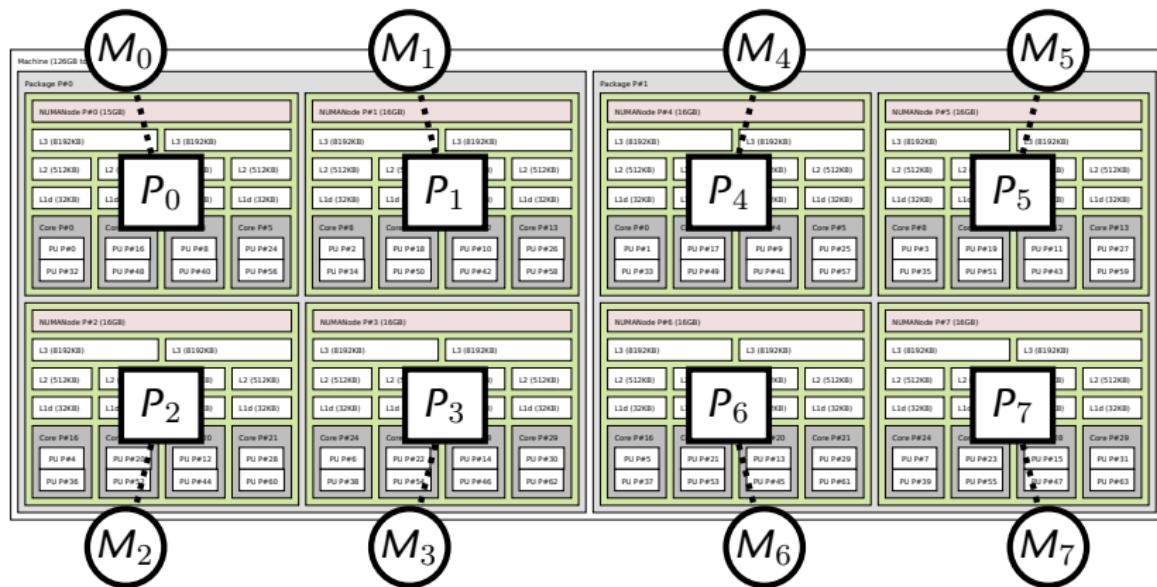
- ▶ SMP node with 2× AMD EYPC “Naples” CPUs
- ▶ 4 × independent memory controller per CPU



# Non-Uniform Memory Access

## Real Example

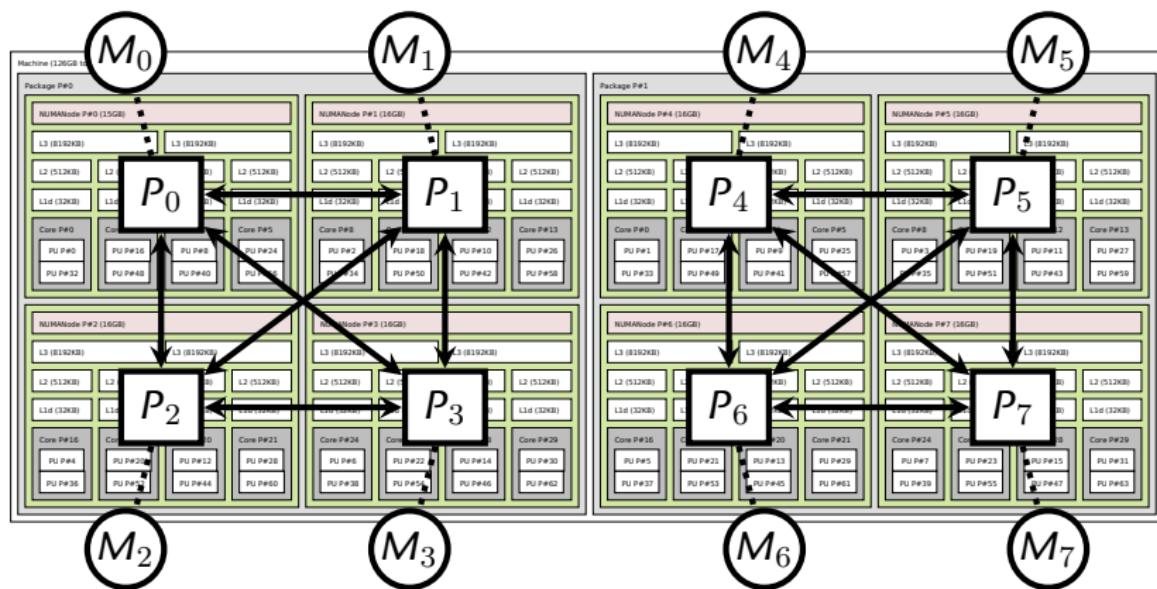
- ▶ SMP node with 2× AMD EYPC “Naples” CPUs
- ▶ 4 × independent memory controller per CPU



# Non-Uniform Memory Access

## Real Example

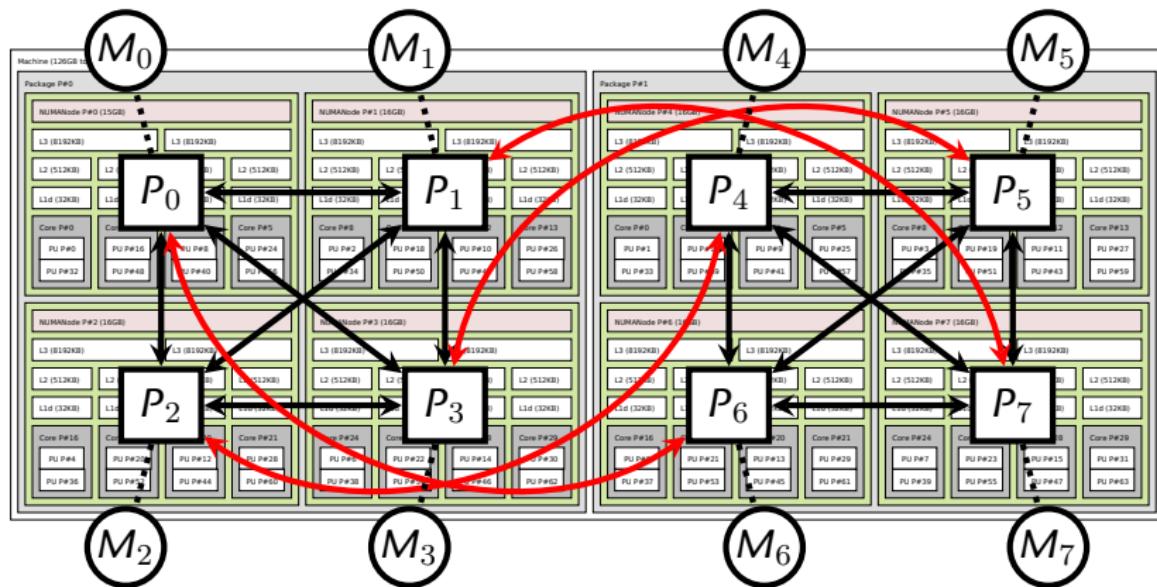
- ▶ SMP node with 2× AMD EYPC “Naples” CPUs
- ▶ 4 × independent memory controller per CPU
- ▶ Complete interconnection inside of a CPU



# Non-Uniform Memory Access

## Real Example

- ▶ SMP node with 2× AMD EYPC “Naples” CPUs
- ▶ 4 × independent memory controller per CPU
- ▶ Complete interconnection inside of a CPU
- ▶ Partial interconnection between the two sockets
- ▶ Red link =  $\approx 2\times$  higher latency



## In practice...

- ▶ machine = nodes connected to a network
  - ▶ MIMD / distributed memory
- ▶ One node = several processors (SMP)
  - ▶ Most likely NUMA
- ▶ Inside a node = all cores access the memory
  - ▶ Shared memory
- ▶ Inside a core = vector instructions
  - ▶ SIMD

Everything, everywhere, all the time.

# Instruction-Level Parallelism (ILP)

```
slwi 10,9,3  
add 8,11,10  
lwzx 10,11,10  
lwz 7,4(8)  
or. 10,10,7  
bne 0,.L146  
addi 5,5,8  
stw 3,0(8)  
stw 4,4(8)  
cmplw 7,6,5  
bne 7,.L24  
lwz 9,144(19)  
li 10,1  
stw 10,20704(31)  
addi 9,9,1
```

code = *ordered sequence of instructions*

## Parallelism inside a core

- ▶ *Pipeline(s)*
- ▶ *Superscalarity*
- ▶ *Out-of-order execution*
- ▶ *Simultaneous Multi-Threading*

⇒ Parallel implementation

⇒ Sequential semantics

# How to Write Parallel Programs?

- ▶ Automatic parallelization

# How to Write Parallel Programs?

- ▶ Automatic parallelization
- ▶ “Annotations” and compiler hints in C programs ([OpenMP](#), [OpenACC](#))

## OpenMP

```
#pragma omp parallel for
for (int i = 0; i < n; i++)
    for (int j = 0; j < n; j++)
        for (int k = 0; k < n; k++)
            C[i][j] += A[i][k] * B[k][j];
```

# How to Write Parallel Programs?

- ▶ Automatic parallelization
- ▶ “Annotations” and compiler hints in C programs ([OpenMP](#), [OpenACC](#))
- ▶ plain C + libraries (pthread, [MPI](#))

## Librairies

```
for(int t=0; t<NUM_THREADS; t++)
    pthread_create(&threads[t], NULL, ThreadFunction, (void *) t);
...
pthread_exit();
```

# How to Write Parallel Programs?

- ▶ Automatic parallelization
- ▶ “Annotations” and compiler hints in C programs ([OpenMP](#), [OpenACC](#))
- ▶ plain C + libraries (pthread, [MPI](#))
- ▶ New languages?

## Go

```
func main() {
    s := []int{7, 2, 8, -9, 4, 0}
    c := make(chan int)
    go sum(s[:len(s)/2], c)
    go sum(s[len(s)/2:], c)
    x, y := <-c, <-c
    fmt.Println(x, y, x+y)
}
```

# Recycling Sequential Programming Languages

## *Single Program Multiple Data (SPMD)*

- ▶ Same sequential code executed many times in parallel
- ▶ Special variable: `rank`
- ▶  $\text{rank} = i$  in the  $i$ -th copy of the program

```
int i = rank;                      /* data parallelism */
for (int j = 0; j < n; j++)
    for (int k = 0; k < n; k++)
        C[i][j] += A[i][k] * B[k][j];

if (rank == 0) {                   /* control parallelism */
    <<Purely sequential section>>
} else {
    <<Wait>>
}
```

# Performance Evaluation

Justice of the peace : wall clock

- ▶  $T_1(n)$ : time needed by the **best** sequential algorithm to process an instance of size  $n$
- ▶  $T_p(n)$ : time needed by the parallel algorithm to process an instance of size  $n$  on  $p$  processors

## Speedup

$$S(n, p) = \frac{T_1(n)}{T_p(n)}$$

## Efficiency

$$E(n, p) = \frac{S(n, p)}{p}$$

# Strong scaling

## Goal when Parallelizing?

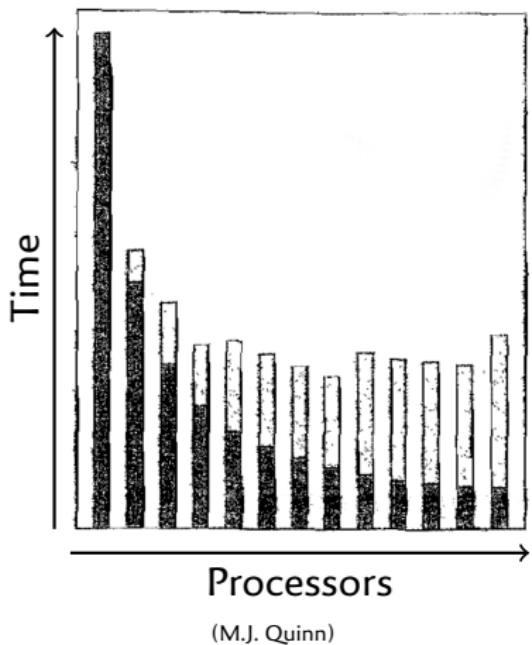
- ▶ **Minimize time-to-solution for a fixed problem**
  - ⇒ Study of the performance with fixed  $n$  and increasing  $p$
- ▶ **Linear speedup** ( $\rightarrow$  ideal): processors are 100% busy, no extra operations

$$S(n, p) = p \quad E(n, p) = 1$$

- ▶ **Sublinear speedup**: processors are less than 100% busy, or extra operations
- ▶ **Superlinear speedup**: hard to envision

# Strong scaling

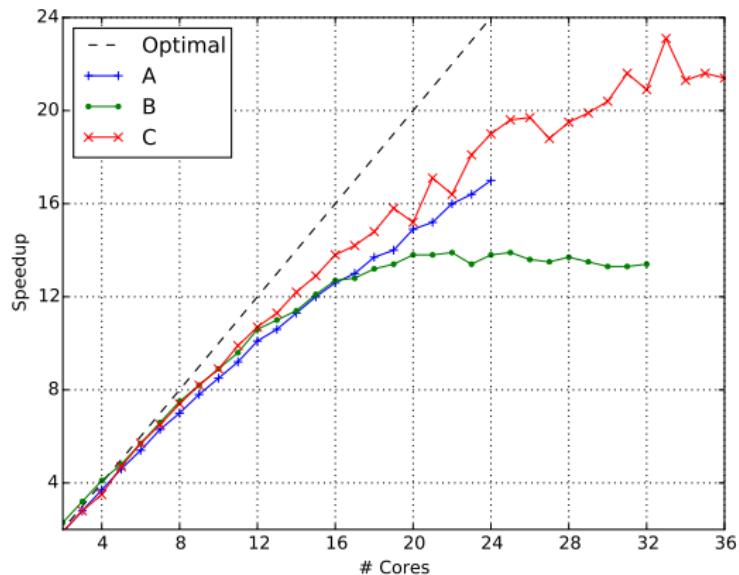
## Real example



- ▶ Gray: computation time
  - ▶ decreases (or stagnate) when  $p \nearrow$
- ▶ White: overhead
  - ▶ Increases when  $p \nearrow$
- ▶ For a given problem size ( $n$ ), there is an optimal number of processors that can be efficiently exploited
  - ▶ Beyond this point, adding extra processors no longer brings any gain

# Real Example (my code)

Compute Rank of Sparse Integer Matrices



| Name | CPU Type        | # CPU | cores/CPU | L3 Cache/CPU |
|------|-----------------|-------|-----------|--------------|
| A    | Xeon E5-2670 v3 | 2x    | 12        | 30MB         |
| B    | Xeon E5-4620    | 4x    | 8         | 16MB         |
| C    | Xeon E5-2695 v4 | 2x    | 18        | 45MB         |

# Amdahl's Law

Sequential Portions Limit Parallel Speedup

Algorithm with a fraction  $f$  of operations that must be performed sequentially. Then

$$S(n, p) \leq \frac{1}{f + (1 - f)/p}$$

$\Rightarrow$  20% of an algorithm is sequential  $\rightsquigarrow$  Speedup  $\leq 5$

Overhead caused by parallelism  $(1 - E(n, p))$  caused by sequential parts and

- ▶ Communications
- ▶ Starting / Synchronizing tasks
- ▶ Load imbalance
- ▶ ...

# Weak scaling

## Goal when Parallelizing?

- ▶ **Solve *bigger* problems with more processors**
- ⇒ Study of the performance when both  $p$  and  $n$  increase

## Amdahl's effect (empirical)

- ▶ When  $n \nearrow$ , computation time often increases *more* than communication time
- ▶ Often allows for greater speedups

## Gustafson-Barsis's Law

A parallel computation is run on  $p$  processors; time spent in sequential operations is measured and represents a fraction  $s$  of the total.

Then :

$$S(n, p) \leq p - (p - 1)s$$

### Example

A parallel computation runs in 100s on 32 processors, including 5s in serial portions. Then  $S(n, p) \leq 30.45$ .