

Introduction to OpenMP


November 14, 2023

Prologue: Thread VS Process


Process : "control flow" + "memory space"

Thread : "control flow"

Specific to each process	Specific to each thread
Memory (Heap)	Instruction pointer
Global variables	Registers
Open Files	Stack (local variables)
child process, signals...	CPU state




Multi-processes




multi-threads

Prologue: Thread VS Process (continued)

Single-thread process:



multi-thread process:



All threads have access to the process' memory space

- ▶ **Shared** variables (between several threads)
 - ▶ *global* variables ("data segment")
 - ▶ Dynamically allocated variables with shared pointer
- ▶ **private** variables ("owned" by a thread)
 - ▶ Local variables (on the stack)
 - ▶ Dynamically allocated variables with private pointer

OpenMP

- 😊 Easier to use than MPI
- 😊 Preserve the original sequential code
- 😊 Code is easier to understand and maintain
- 😊 Allows progressive parallelization
- 🤪 **Shared memory** machines only (e.g. ONE server)
- 🤪 Works better on specific code patterns (loop nests, ...)

MPI


- 😊 Designed for **distributed memory** machines
- 😊 Also works fine on shared-memory machines
- 😊 Separated memory spaces (no conflicts)
- 🤪 Algorithmic modifications often necessary
- 🤪 Harder to use
- 🤪 Performance depends on the network

OpenMP: timeline

- ▶ 1997, a consortium of industrials and academics adopt OpenMP (*Open Multi Processing*) as a **standard**. Fortran, C and C++ interface
- ▶ Version 2.5 (2000, gcc 4.2) : *lean and mean* (for loops)
- ▶ Version 3.0 (2008, gcc 4.4) : new *task* concept
- ▶ Version 3.1 (2008, gcc 4.7) : better tasks, `atomic`
- ▶ Version 4.0 (2013, gcc 4.9) : SIMD and *devices* (GPU, ...)
- ▶ Version 4.5 (2013, gcc 6) : better SIMD, more GPU, ...
- ▶ Version 5.1 (2020, gcc 10) : `atomic compare + ...`
 - ▶ v5.1 of the spec is much harder to read than v4.5
- ▶ Version 5.2 (2021) : clean up the spec...
- ▶ Version 6.0 (2024) : ???

Principle

- ▶ A single process runs on a machine.
- ▶ The corresponding thread is the “**master** thread” (number 0)
- ▶ It occasionally spawns other threads to run parallel computation, then waits for them to complete (*fork and join* model)



- ▶ Declaring parallel sections in the code is done using **OpenMP directives**
- ▶ Specific **memory model** (shared / private variables)

Using OpenMP

- ▶ **Compile-time** directives (#pragma in C)
 - ▶ Interpreted by OpenMP-aware compilers
 - ▶ **Silently** ignored otherwise
 - ▶ Tell the compiler to generate parallel code
 - ▶ gcc: -fopenmp option
- ▶ Must **link** against OpenMP run-time library
 - ▶ gcc: -fopenmp option
- ▶ At **run-time**: **environment variables** allow some control over OpenMP

OpenMP in One Slide

```
// sequential prologue
#pragma omp parallel for
for (int i = 0 ; i < n ; i++) {
    /*
     * All the iterations of this loop can
     * be executed in parallel
     */
}
// sequential epilogue
```

gcc -fopenmp prog_omp.c -o prog_omp

- ▶ “normal” sequential program until `#pragma omp`
- ▶ A **team of threads** is created
- ▶ The iterations of the loop are **distributed** between them
- ▶ **Barrier** at the end of the loop
- ▶ “normal” sequential program then

Conditional Compilation and Run-Time OpenMP Functions

Conditional compilation

```
#ifdef _OPENMP  
    // Code included only if the compiler supports OpenMP  
    // With gcc, only if the -fopenmp option has been given  
#endif
```

OpenMP run-time functions

With `#include <omp.h>`

- ▶ Enable a SPMD style of programming (as in MPI)
- ▶ `omp_get_num_threads()`
- ▶ `omp_get_thread_num()`
- ▶ `omp_set_num_threads()`
- ▶ ...

Hello world

Program

```
#ifdef _OPENMP
#include <omp.h>
#endif

int main()
{
    #pragma omp parallel
    {
        #ifdef _OPENMP
        printf("Hello world, thread %d/%d\n",
            omp_get_thread_num(),
            omp_get_num_threads());
        #else
        printf("Hello world\n");
        #endif
    }
}
```

Execution

```
$ gcc hello.c -o hello
$ ./hello
Hello world
$ gcc -fopenmp hello.c \
  -o hello
$ export OMP_NUM_THREADS=4
$ ./hello
Hello world, thread 0/4
Hello world, thread 3/4
Hello world, thread 1/4
Hello world, thread 2/4
```

OpenMP Directives

```
#pragma omp directive [clause[[, ]clause]...]
```

Barrier (synchronization) at the end by default

Using directives

- ▶ Jumping out of a parallel section is forbidden (`goto`, `setjmp/longjmp`, ...)
- ▶ One directive per `#pragma omp`
- ▶ Case-sensitive
- ▶ Directives $\subseteq \{ \text{parallel, for, sections, section, single, master, critical, barrier, atomic, flush, ordered, threadprivate, ...} \}$


OpenMP Memory Model

Reminder

Variable : identifier denoting a memory address

Variables present in the original sequential code can be declared as *shared* or *private* with OpenMP.

- ▶ **Shared**: all threads access the memory address of the original variable
- ▶ **Private**: each thread "owns" a **copy** of the original variable (all located at different memory addresses)



(after *An Overview of OpenMP 3.0*, R. van der Pas)

OpenMP Memory Model (continued)

- ▶ Variables declared **before** a parallel region are **shared by default**
- ▶ Their status can be modified using **clauses** in OpenMP directives
 - ▶ `private`, `shared`, `firstprivate`, `lastprivate`,
`default(shared)`, `default(none)`, `reduction`, `copyin`
- ▶ Local variables of each thread are private (cf. infra)

parallel Directive

```
#pragma omp parallel [clause[, ]clause]...]  
structured block
```

- ▶ Assemble a **thread team** (creation/recycling)
- ▶ **All** threads run the *structured block*.

Associated clauses

- ▶ `if(cond): cond == False` → no threads
 - ▶ E.g., don't use all of the machine for small problem
- ▶ `private (var_list), firstprivate (var_list)`
- ▶ `reduction` (cf. infra)
- ▶ `num_threads(int)`: force the size of the thread team

```

int main()
{
    ...
    initialization();
    #pragma omp parallel ...
    {
        parallel_computation();
    }
    post_processing();
}

```

How to choose the number of threads

By decreasing order of priority:

Compile-time	: #pragma omp parallel num_threads(16)
Run-time	: <i>omp_set_num_threads(4)</i>
Environment variable	: export OMP_NUM_THREADS=4

Predefined Variables are Shared by Default

Program

```
#include <omp.h>
#include <stdio.h>

int main()
{
    int c=0;

    #pragma omp parallel
    {
        c++;
        printf("c=%d thread %d\n",
               c, omp_get_thread_num());
    }
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./a.out
c=1 thread 3
c=2 thread 0
c=3 thread 1
c=4 thread 2
```


Beware of Conflicts!

Program

```
#include <omp.h>
#include <stdio.h>

int main()
{
    int c = 0;

    #pragma omp parallel
    {
        for (int i=0; i<100000; i++)
            c++;
        printf("c=%d thread %d\n",
            c, omp_get_thread_num());
    }
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./a.out
c=100000 thread 0
c=200000 thread 3
c=270620 thread 2
c=286162 thread 1
```

private Clause

private variable:

- ▶ Each thread owns a (private) local copy
- ▶ **Not initialized**

BUG :

Program

```
int main()
{
    int a = 100;
    #pragma omp parallel private(a)
    {
        /* This "a" is not the same
           as before */
        a = a + 10;
        printf("a=%d\n", a);
    }
    printf("After a=%d\n", a);
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./a.out
a=-1208433038
a=-22
a=-22
a=-22
After a=100
```

firstprivate Clause

firstprivate Variables;

- ▶ Each thread owns a (private) local copy
- ▶ Initialized with the preexisting value

Program

```
int main()
{
    int a = 100;
    #pragma omp parallel \
        firstprivate(a)
    {
        a = a + 10;           // idem...
        printf("a=%d\n", a);
    }
    printf("After a=%d\n", a);
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./a.out
a=110
a=110
a=110
a=110
After a=100
```

Local Variables

- ▶ All local variables in functions called from a parallel section are owned by the corresponding thread (on their stacks)
- ▶ Same goes for local variables declared inside the *structured block*

Program

```
void func()
{
    int a = 10;
    a += omp_get_thread_num();
    printf("a=%d\n", a);
}

int main()
{
    #pragma omp parallel
    func();
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./test2
10
11
12
13
```

Local Variables: My Opinion

Complex (avoid)

```
int main()
{
    int a;
    #pragma omp parallel private(a)
    {
        a = ...;
        ...
    }
}

/*****/

int main()
{
    int a;
    #pragma omp parallel firstprivate(a)
    {
        ... a ...
    }
}
```

Simple (better)

```
int main()
{
    #pragma omp parallel
    {
        int a = ...;
        ...
    }
}

/*****/

int main()
{
    int a;
    #pragma omp parallel
    {
        int b = a;
        ... b ...
    }
}
```

Reminders / Details

```
#pragma omp parallel
```

- ▶ A **thread team** is created
- ▶ The **encountering thread** belongs to it (**master**)
- ▶ All threads of the team execute the *structured block*
- ▶ Identifiers ("ranks") 0 (master), 1, 2, ..., #threads - 1
 - ▶ `omp_get_num_threads()`, `omp_get_thread_num()`
- ▶ **Barrier** at the end
- ▶ Encountering thread then resumes sequential execution

for Directive

```
#pragma omp for [clause[[, ]clause]...]  
⟨ for loop ⟩
```

Worksharing directive

Threads of the team cooperate and **divide** the work between them

Associated clauses :

- ▶ `private (variable_list), firstprivate (variable_list), lastprivate (variable_list)`
- ▶ `reduction(operator: variable_list)`
- ▶ `ordered`
- ▶ `collapse(n)`
- ▶ `schedule(type, size)`
- ▶ `nowait`

for Directive (continued)

Canonical loop form

`for (init_expr ; cond ; increment)`

- ▶ Integer iteration variable
- ▶ Loop counter updated with `++`, `--`, `+=`, `-=`,
`var=var+inc`, `var=inc+var`, `var=var-inc`
- ▶ integer increment
- ▶ Condition: `<`, `>`, `<=`, `>=`. Bound is a fixed expression
- ▶ No early exit (`break`, `return`, `exit`)
- ▶ `continue` is allowed

Consequences of the `for` directive:

- ▶ Implicit barrier at the end of the loop (except if `nowait`)
- ▶ No barrier at the beginning
- ▶ **The iteration variable is private**

Example

```
int main()
{
    int t[100];
    #pragma omp parallel
    {
        #pragma omp for
        for (int i = 0; i < 100; i++)
            t[i] = i;
    }
}
```

With 4 threads, the first one may **for instance** compute the $t[i]$ from 0 to 24, the second from 25 to 49, ...

Short Form of the for directive

```
#pragma omp parallel for [clause[[, ]clause]...]
for loop
```

Admits all clauses of `parallel` and `for`, except `nowait`.

```
#pragma omp parallel
#pragma omp for
for (int i = 0 ; i < n ; i++) {
    ....
}

#pragma omp parallel for
for (int i = 0 ; i < n ; i++) {
    ....
}
```

The reduction Clause

Program

```
int main()
{
    int a[4][4], s=0;
    for (int i = 0; i < 4; i++)
        for (int j = 0; j < 4; j++)
            a[i][j] = i * 4 + j;
    #pragma omp parallel for reduction(+:s)
    for (int i = 0 ; i < 4 ; i++) {
        for (int j = 0; j < 4; j++)
            s += a[i][j];
        printf("PAR=%d : i=%3d s=%d\n",
            omp_get_thread_num(), i, s);
    }
    printf("SEQ s=%d\n", s);
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./a.out
PAR=0 : i= 0 s=6
PAR=1 : i= 1 s=22
PAR=2 : i= 2 s=38
PAR=3 : i= 3 s=54
SEQ s=120
```

The reduction Clause (continued)

- ▶ Operators : +, -, * , &, |, ^, &&, ||, min, max
- ▶ Can define your own

Authorized on `omp for`, `omp parallel`, ...

```
int main()
{
    int m = 0;
    #pragma omp parallel reduction(max:m)
    {
        int tid = omp_get_thread_num();
        m = f(tid);
    }
    ...
}
```

New Features in reduction



- ▶ reduction on *arrays*
- ▶ ... or array *slices*
- ▶ Starting with OpenMP 4.5
 - ▶ November 2015, gcc \geq 6.1

```
double *A = malloc(n * sizeof(*A));
...
#pragma omp parallel reduction(+:A[0:n])
{
    // Each thread own its own copy of A[0:n]
    for (int i = 0; i < n; i++) {
        A[i] = ....
    }
    ....
}
// A[i] contains the sum of the
// private A[i]'s of all threads
```

Load Balancing in Loops

`omp for` admits load-balancing clauses: `schedule` and `nowait`

- ▶ `nowait` clause:
Removes the automatic barrier at the end of the loop
- ▶ `schedule(mode, chunk_size)` clause:
4 modes: `static`, `dynamic`, `guided`, `runtime`

Implementation-dependent if not specified (`static` on `gcc`)

schedule Example

```
#define MAX 10
int main()
{
    int a[MAX];
    #pragma omp parallel
    {
        int imax = 0;
        int imin = MAX;
        #pragma omp for schedule(static)
        for (int i = 0; i < MAX; i++) {
            imin = (i < imin) ? i : imin;
            imax = (i > imax) ? i : imax;
            a[i] = 1;
            sleep(1); /* simulate computation */
            printf("%3d:%3d\n", omp_get_thread_num(), i);
        }
        printf("T%d imin=%d imax=%d\n", omp_get_thread_num(), imin, imax);
    }
}
```

static and dynamic Clauses

schedule(...)

static

```
1: 3
3: 8
2: 6
0: 0
1: 4
3: 9
2: 7
0: 1
1: 5
0: 2
```

```
T1 imin=3 imax=5
T0 imin=0 imax=2
T2 imin=6 imax=7
T3 imin=8 imax=9
```

static, 2

```
1: 2
3: 6
2: 4
0: 0
1: 3
3: 7
0: 1
2: 5
0: 8
0: 9
```

```
T0 imin=0 imax=9
T1 imin=2 imax=3
T3 imin=6 imax=7
T2 imin=4 imax=5
```

dynamic, 2

```
1: 4
0: 0
2: 6
3: 2
1: 5
2: 7
0: 1
3: 3
1: 8
1: 9
```

```
T1 imin=4 imax=9
T3 imin=2 imax=3
T2 imin=6 imax=7
T0 imin=0 imax=1
```


Schedule

- ▶ `schedule(static)`: block distribution
 - ▶ `schedule(static, n)`: block-cyclic distribution.
 - ▶ block size n
 - ▶ `schedule(dynamic, n)`: chunks of n iterations are affected to available threads ($n = 1$ by default).
 - ▶ `guided`: like `dynamic` but the chunk size is proportional to the number of remaining iterations
 - ▶ `auto`: the OpenMP runtime does its magic
 - ▶ `runtime`: choice is deferred until runtime
- Example : `export OMP_SCHEDULE="static,1"`

single Directive

```
#pragma omp single directive [clause[[, ]clause]...]
structured block
```

Goal

Sequential portion inside a parallel region, *i.e.* a code chunk executed by a **single** thread

- ▶ It can be by *any* thread
- ▶ Implicit barrier at the end of `single`
- ▶ `nowait` and `copyprivate` are incompatible clauses

Possible clauses:

- ▶ `private` (*variable_list*), `firstprivate` (*variable_list*),
`copyprivate` (*variable_list*)
- ▶ `nowait`

single Example: Prefix-Sum

```
int s = 0;
for (int i = 0; i < n; i++) {
    int tmp = A[i];
    A[i] = s;
    s += tmp;
}
```

```
#pragma omp parallel
{
    int t = omp_get_thread_num();
    S[t] = 0;
    #pragma omp for schedule(static)
    for (int i = 0; i < n; i++)
        S[t] += A[i];
    #pragma omp single
    {
        int p = omp_get_num_threads();
        int s = 0;
        for (int i = 0; i < p; i++) {
            int tmp = S[i];
            S[i] = s;
            s += tmp;
        }
        int s = S[t];
        #pragma omp for schedule(static)
        for (int i = 0; i < n; i++) {
            int tmp = A[i];
            A[i] = s;
            s += tmp;
        }
    }
}
```

Synchronizations in OpenMP

Several options:

- ▶ Barriers
- ▶ `atomic` and `critical`
- ▶ Locks via OpenMP runtime functions (not covered):

`omp_init_lock()`

`omp_{set,test}_lock()`

`omp_unset_lock()`

`omp_destroy_lock()`

barrier Directive

```
#pragma barrier
```

All threads must *enter* the barrier before any of the threads continue execution beyond the barrier

Problems with C syntax

```
if (n != 0)
    #pragma omp barrier    // syntactically incorrect
```

```
if (n != 0) {
    #pragma omp barrier    // OK
}
```

critical Directive

```
#pragma omp critical [name]  
structured block
```

- ▶ Only a single thread may enter the *structured block* simultaneously
- ▶ “**critical section**”
- ▶ An incoming thread has to wait while another thread executes the *structured block*
- ▶ The *name* allows to distinguish distinct critical sections.

critical Example

Adding an Item to a Linked List

```
struct item_t {  
    void *data;  
    struct item_t *next;  
};  
struct item_t *list;           /* global variable */  
  
void append(void *payload)  
{  
    struct item_t *new_item = malloc(sizeof(*new_item));  
    new_item->data = payload;  
    #pragma omp critical  
    {  
        new_item->next = list;  
        list = new_item;  
    }  
}
```

atomic Directive

```
#pragma omp atomic [ read | write | update ]  
atomic-update
```

- ▶ The *atomic-update* is **atomic** (cannot be interrupted)
- ▶ *atomic-update* of the type:
 - ▶ `v = x; OR x = expr;`
 - ▶ `x += expr; OR x = x + expr;`
 - ▶ Works with `+`, `*`, `-`, `/`, `&`, `&&`, `^`, `|`, `||`, `>>`, `<<`
 - ▶ *expr* must not reference *x*.
- ▶ Only reading/writing/updating *x* is atomic.
- ▶ Evaluation of *expr* is not

Program

```
#include <omp.h>

int main()
{
    int c = 0;
    #pragma omp parallel
    {
        for (int i = 0; i < 100000; i++) {
            #pragma omp atomic
            c++;
        }
        printf("c=%d thread %d\n",
            c, omp_get_thread_num());
    }
}
```

Execution

```
$ export OMP_NUM_THREADS=4
$ ./a.out
c=100000 thread 0
c=294308 thread 2
c=394308 thread 3
c=400000 thread 1
```

atomic capture

Since OpenMP 3.0 (2008)

```
#pragma omp atomic capture  
structured block
```

- ▶ Capture = saving the previous value ("capture") + update
- ▶ With the capture clause, *structured block* can be:

```
{v = x; x = expr;}
```

```
{v = x; x += expr;}
```

```
{x += expr; v = x;}
```

```
{v = x; x = x + expr;}
```

```
{v = x; x = expr + x;}
```

```
{x = x + expr; v = x;}
```

```
{x = expr + x; v = x;}
```

```
{v = x; x++;}
```

```
{v = x; ++x;}
```

```
{++x; v = x;}
```

```
{x++; v = x;}
```

atomic capture Example

Appending an Item to a Preallocated Array

```
void *A[];  
int n = 0;           /* #items in the array */  
  
void atomic_append(void *payload)  
{  
    int i;  
    #pragma omp atomic capture  
    {  
        i = n;        /* equivalent to: i = n++ */  
        n += 1;  
    }  
    A[i] = payload;  
}
```

atomic compare

Since OpenMP 5.1 (needs gcc 12.2, August 2022!)

"Compare-and-Swap"

```
#pragma omp atomic compare  
if (x == old) x = new;
```

"Compare-and-Swap" with capture

```
#pragma omp atomic compare capture  
{ if (x == old) { x = new; }; else { cap = x; } }
```

"Compare-and-Swap", most general form

```
#pragma omp atomic compare capture  
{ ok = x==old; if (ok) { x = new; } else { cap = x; } }
```

- All CPUs have *hardware* support for this operation

atomic compare Example

Adding an Item to a Linked List

```
struct item_t {  
    void *data;  
    struct item_t *next;  
}  
struct item_t *list;                                /* global variable */  
  
void atomic_append(void *payload)  
{  
    struct item_t *new_item = malloc(sizeof(*new_item));  
    new_item->data = payload;  
    bool ok = 0;  
    while (!ok) {  
        new_item->next = list;  
        #pragma omp atomic compare capture  
        {  
            ok = (list == new_item->next);  
            if (ok)  
                list = new_item;  
        }  
    }  
}
```

"Transactions" Using *Compare-And-Swap*

It is possible to do almost **anything** using *Compare-And-Swap*!

Idiom: *Compare-And-Swap Loop*

1. [Begin.] $x_{old} \leftarrow x$
2. [Work.] Compute an updated x_{new}
 - ▶ This may read x ...
 - ▶ ... and x may be modified by another thread
3. [Commit.] `ok = (x == x_old); if (ok) x = x_new;`
4. [Repeat.] If not OK, go back to 1.

atomic compare Example

Insertion Into a Hash Table With Linear Probing

Hash Table With Linear Probing

```
void insert(void *H, void *item)
{
    int i = hash_function(item);           // hash
    while (H[i] != EMPTY)                  // locate empty slot
        i = (i + 1) % HASHTABLE_SIZE;
    H[i] = item;                           // insert
}
```

Thread-safe version

```
void atomic_insert(void *H, void *item)
{
    int i = hash_function(item);
    book ok = 0;
    while (!ok) {
        #pragma omp atomic capture
        { ok = H[i] == EMPTY; if (ok) H[i] = item; }
        i = (i + 1) % HASHTABLE_SIZE;
    }
}
```

Differences Between atomic and critical

- `atomic:`
- ▶ Intended to update variables
 - ▶ Depends on hardware support
 - ▶ “Atomic” CPU instructions (compare-and-swap or ll/sc)
 - ▶ Implemented with locks as a last resort
 - ▶ Lighter overhead *a priori*
 - ▶ To be preferred if possible

- `critical:`
- ▶ Intended to encompass a larger portion of code
 - ▶ Implemented using locks
 - ▶ Heavier overhead

atomic VS critical: Tricky Example

Appending an Item to a Dynamic Array **with resizing**

```
void *A[];
int n = 0;                                /* #items inserted in A[] */
size = 0;                                  /* allocated size of A */

void append(void *payload)
{
    #pragma omp critical
    {
        if (n >= size) {                  /* array too small? Enlarge it */
            size = 2 * size + 1;
            A = realloc(A, size);          /* change A */
        }
        A[n] = payload;                    /* finally, insert */
        n += 1;
    }
}
```

Without locks / critical?

- ▶ Possible, but quite difficult
- ▶ Nice challenge! Try it at home!

Lock-Free Algorithms

Definition

A function is **lock-free** if, when it is called by several threads at the same time, at least one of the invocations finishes in a finite number of steps (whatever the others do, even if they block).

⇒ if a function acquires a lock or enters a critical section, it cannot be *lock-free*.

Challenge

Design a *lock-free* data structure to hold a set of N integers. API:

```
void setup(int N);           /* initialise  $S = \{0, 1, \dots, N-1\}$  */
bool remove(int i);          /* Remove  $i$  from  $S$ . Return 1 if  $i$  was in  $S$  */
int extract_min();            /* let  $k = \min(S)$  in  $\text{remove}(k)$  */
```

Golden rule of multi-thread programming

ALL potentially conflicting accesses* to **shared** variables **MUST** be protected (atomic, critical, ...).

* at least one of them is a write

The Same Example Again

Program

```
#include <stdio.h>


int main()
{
    int c = 0;
    #pragma omp parallel for
    for (int i = 0; i < 100000; i++) {
        c++;
    }
    printf("c=%d\n", c);
}
```

Execution

```
$ ./a.out
c=15074
```


Interleaved Increments

Load-Store Hardware Architectures



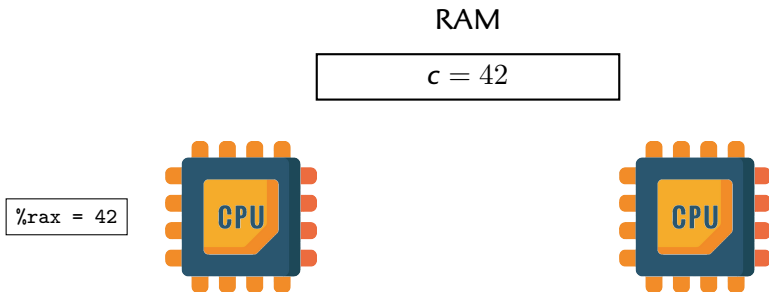
Interleaved Increments

Load-Store Hardware Architectures



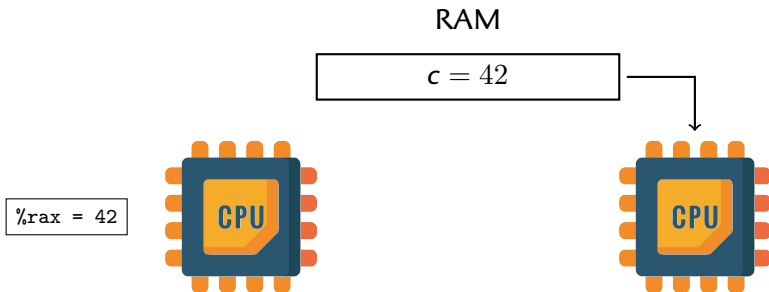
Interleaved Increments

Load-Store Hardware Architectures



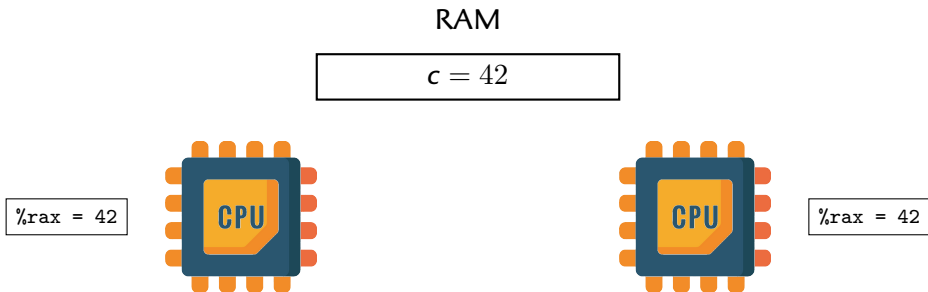
Interleaved Increments

Load-Store Hardware Architectures



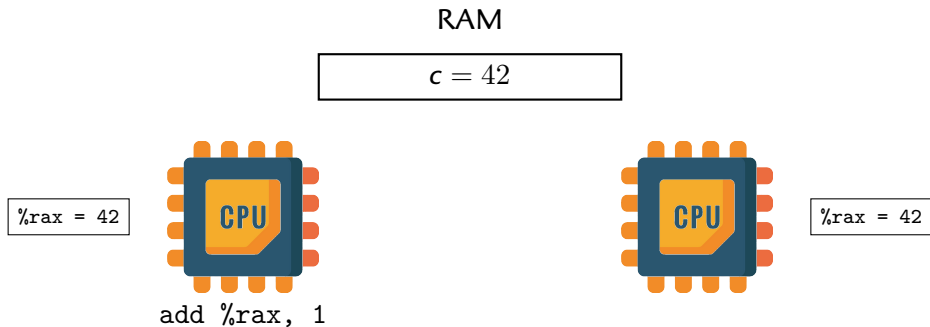
Interleaved Increments

Load-Store Hardware Architectures



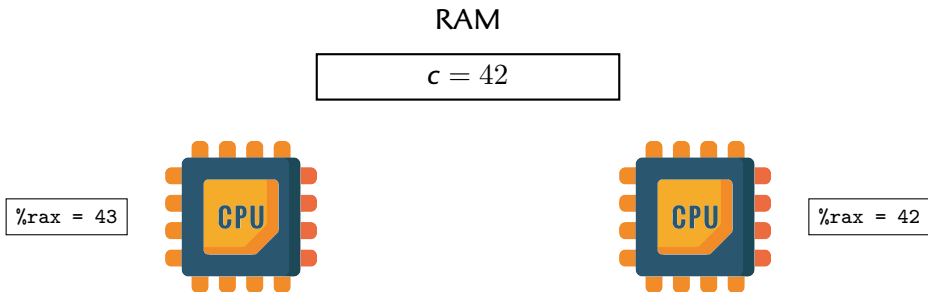
Interleaved Increments

Load-Store Hardware Architectures




Interleaved Increments

Load-Store Hardware Architectures



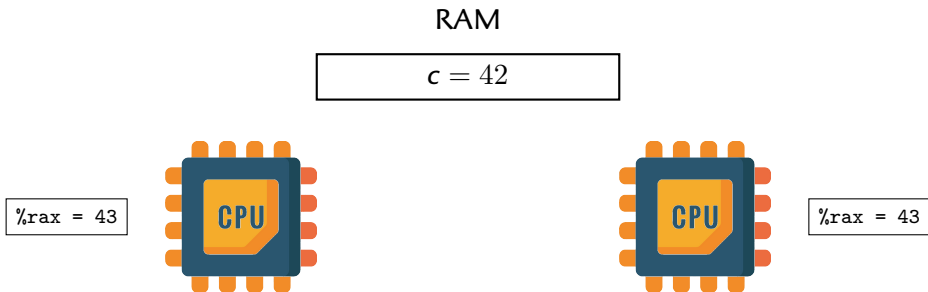
Interleaved Increments

Load-Store Hardware Architectures



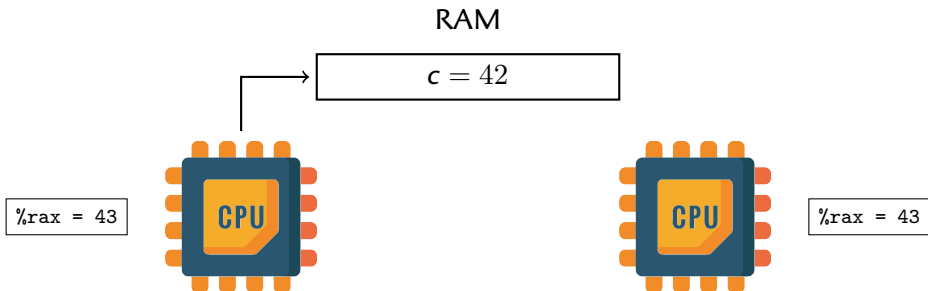
Interleaved Increments

Load-Store Hardware Architectures




Interleaved Increments

Load-Store Hardware Architectures




Interleaved Increments

Load-Store Hardware Architectures




Interleaved Increments

Load-Store Hardware Architectures



Interleaved Increments

Load-Store Hardware Architectures





Bad

```
#include <stdio.h>

int main()
{
    int c = 0;
    #pragma omp parallel for
    for (int i = 0; i < 100000; i++) {
        c++;
    }
    printf("c=%d\n", c);
}
```



Good

```
#include <stdio.h>

int main()
{
    int c = 0;
    #pragma omp parallel for
    for (int i = 0; i < 100000; i++) {
        #pragma omp atomic update
        c++;
    }
    printf("c=%d\n", c);
}
```

Load-store architectures:

- ▶ $x++$, $x += 10$, etc. are **not atomic**
- ▶ Not even when performed by a **single CPU instruction**

Tasks

Since OpenMP 3.0 (2008)

French speakers: point d'orthographe

Tache marque, salissure, souillure (*stain*). *Une tache d'encre*

Tâche travail à exécuter (*task*). *Une tâche ardue*

Tasks

Since OpenMP 3.0 (2008)

French speakers: point d'orthographe

Tache marque, salissure, souillure (*stain*). *Une tache d'encre*

Tâche travail à exécuter (*task*). *Une tâche ardue*

Reminder: `#pragma omp parallel`

- ▶ A **thread team** is assembled
- ▶ ...
- ▶ An **implicit task** is scheduled on each thread
- ▶ It is **tied** (cannot migrate to another thread)

task = code + associated data ("closure"), executed by a thread

omp task Directive

```
#pragma omp task [clause], [clause], ...  
structured block
```


- ▶ The encountering thread creates an **explicit task**
- ▶ Execution can be **immediate** or **deferred**
- ▶ Execution by one of the threads of the team... if available

omp task Directive

```
#pragma omp task [clause], [clause], ...  
structured block
```

- ▶ The encountering thread creates an **explicit task**
- ▶ Execution can be **immediate** or **deferred**
- ▶ Execution by one of the threads of the team... if available
- ▶ A thread may suspend a task and start/resume another one...
- ▶ ... but only at a *task scheduling point*

"Task Pool"



Developer specifies tasks in application
Run-time system executes tasks

(An Overview of OpenMP 3.0, R. van der Pas,
IWOMP2009)

Synchronizing Tasks

Fork/Join Model

"Normal" barrier

- ▶ Implicit: at the end of a parallel region, of `omp for, ...`
- ▶ Explicit: **`#pragma omp barrier`**

Guarantee: All tasks created by threads in the team are completed after the barrier

Task barrier: **`#pragma omp taskwait`**

- ▶ The current task waits for completion of its **child** tasks
- ▶ Only direct children, not descendant

Example: Iterating Over a Linked List

```
struct item_t {  
    void *data;  
    struct item_t *next;  
};  
struct item_t *list;
```

Sequential

```
struct item_t *e = list;  
  
while (e != NULL) {  
    process(e->data);  
    e = e->next;  
}
```

With tasks

```
struct item_t *e = list;  
#pragma omp parallel  
#pragma omp single  
while (e != NULL) {  
    #pragma omp task  
    process(e->data);  
    e = e->next;  
}
```

Example: Visiting the Nodes of a Binary Tree

```
struct tree_t {  
    ...  
    struct tree_t *left, *right;  
};  
struct tree_t *root;
```

Sequential

```
void walk(struct tree_t *t)  
{  
    ...  
    if (t->left)  
        walk(t->left);  
    if (t->right)  
        walk(t->right);  
}  
  
walk(root);
```

With tasks

```
void walk(struct tree_t *t)  
{  
    ...  
    if (t->left)  
        #pragma omp task  
        walk(t->left);  
    if (t->right)  
        #pragma omp task  
        walk(t->right);  
}  
#pragma omp parallel  
#pragma omp single  
walk(root);
```

omp task Directive (continued)

```
#pragma omp task [clause], [clause], ...  
structured block
```

Associated clauses:

- ▶ `private (variable_list)`, `firstprivate (variable_list)`, `shared (variable_list)`
- ▶ `default(shared | none)`
- ▶ `untied`
- ▶ `depend(dependance-type: list)`
- ▶ `if(expression)`
- ▶ ...

Tasks: Variable Scoping

- ▶ Most useful data attribute with tasks: **firstprivate**
- ~> By default on all variables...
- ▶ ...**except** if they are already **shared**
 - ▶ Global variables
 - ▶ Declared before the parallel section
 - ▶ Explicitly tagged as shared

Be careful with shared variables on the stack

```
void f()
{
    int i = 3;
    #pragma omp task shared(i)
    printf("%d\n", i);
}

#pragma omp parallel
#pragma omp single
f();
```

Tasks: When is shared Necessary?

```
struct tree_t {
    ...
    struct tree_t *left, *right;
};
struct tree_t *root;

/* Return \# nodes in the tree. */
int size(struct tree_t *t)
{
    int s_left = 0, s_right = 0;
    if (t->left)
        #pragma omp task shared(s_left)
        s_left = size(t->left);
    if (t->right)
        #pragma omp task shared(s_right)
        s_right = size(t->right);
    #pragma omp taskwait
    return 1 + s_left + s_right;
}

#pragma omp parallel
#pragma omp single
printf("%d\n", size(root));
```

Tasks: granularity

Creating a task has a non-trivial cost

Don't create microscopic tasks

- ▶ *if* Clause from the `omp task` directive
 - ▶ E.g. `#pragma omp task if(prof < PROF_MAX)`
 - ▶ The task is **created anyway...**
 - ▶ ...but executed immediately by the encountering thread
- ▶ *if* instruction:


```
if (prof < PROF_MAX) {  
    #pragma omp task  
    stuff(...);  
} else {  
    stuff(...);  
}
```

→ to be preferred

Tasks: dependances

Example (Christian Terboven)

```
void blocked_cholesky( int NB, float A[NB][NB] ) {  
    int i, j, k;  
    for (k=0; k<NB; k++) {  
        #pragma omp task depend(inout:A[k][k])  
        spotrf (A[k][k]) ;  
        for (i=k+1; i<NT; i++)  
            #pragma omp task depend(in:A[k][k]) depend(inout:A[k][i])  
            strsm (A[k][k], A[k][i]);  
        // update trailing submatrix  
        for (i=k+1; i<NT; i++) {  
            for (j=k+1; j<i; j++)  
                #pragma omp task depend(in:A[k][i],A[k][j])  
                depend(inout:A[j][i])  
                sgemm( A[k][i], A[k][j], A[j][i]);  
            #pragma omp task depend(in:A[k][i]) depend(inout:A[i][i])  
            ssyrk (A[k][i], A[i][i]);  
        }  
    }  
}
```



* image from BSC

Nested Parallelism

- ▶ A parallel directive inside a parallel directive
- ▶ In most cases, this is a no-op.
 - ▶ No more “available” threads to work on the nested parallel region: they are all in the parent parallel region.
- ▶ It is possible to force the creation of a new thread teams (cf. spec)

My opinion

- ▶ **STRICTLY USELESS** to do it voluntarily
- ▶ Could happen if a parallel program invoked a parallel library inside a parallel region

OpenMP Runtime Library

Some common functions (with `#include <omp.h>`):

- ▶ `void omp_set_num_threads(int num_thread)` : sets the number of threads to use in subsequent parallel regions
- ▶ `int omp_get_num_threads(void)` : return the number of threads in the current team
- ▶ `int omp_get_max_threads(void)` : upper bound on #threads in new teams
- ▶ `int omp_get_thread_num(void)`
- ▶ `int omp_get_num_procs(void)`
- ▶ `int omp_in_parallel(void)` : return true if we are inside a parallel region
- ▶ `double omp_get_wtime(void)`

Environment Variables

- ▶ OMP_NUM_THREADS
- ▶ OMP_SCHEDULE :
 export OMP_SCHEDULE="static,4"
 export OMP_SCHEDULE="dynamic"
- ▶ ...

Cheat Sheet

Most common

- ▶ The *must*: `#pragma omp parallel` for
- ▶ `atomic` and `critical` directives
- ▶ `schedule` and `reduction` clauses

Less common: SPMD mode

- ▶ Parallel region with explicit `#pragma omp for` inside
- ▶ `omp_get_thread_num()` and `omp_get_num_threads()`
- ▶ `barrier` or `single` directives

Rare (reserved for the most delicate cases...)

Everything else!